

AI Forensic Lab - Image Forensics Benchmark Report

Run ID: run_20260808_222229_hf_vit_deepfake | Primary Model: PyTorch ViT DeepFake Classifier
Generated: 2026-08-10 17:24:02 UTC

1. Executive Summary

Evaluated 3 benchmark images (1 Human, 2 AI) using 'PyTorch ViT DeepFake Classifier'. Overall Accuracy reached 33.33% with F1-Score of 0.0000 and average inference latency of 27347.7 ms.

2. Confusion Matrix

Actual / Predicted	Predicted Human	Predicted AI
Actual Human	TN = 1 (100.0%)	FP = 0 (0.0%)
Actual AI	FN = 2 (100.0%)	TP = 0 (0.0%)

3. Key Evaluation Metrics

Metric	Value	Description
Accuracy	33.33%	Overall correct classification rate
Precision	0.0000	Proportion of AI predictions that were AI
Recall (Sensitivity)	0.0000	Proportion of actual AI images detected
Specificity	1.0000	Proportion of actual Human images detected
F1 Score	0.0000	Harmonic mean of Precision and Recall
False Positive Rate	0.00%	Authentic human images misclassified as AI
False Negative Rate	100.00%	AI-generated images missed
ROC-AUC	0.0	Area under ROC curve across thresholds

4. Evaluation Warnings & Recommendations

- [WARNING] Accuracy may not represent real-world performance because the benchmark dataset is imbalanced (Human: 33.3333%, AI: 66.6667%).
- [RECOMMENDATION] AI-generated images are frequently missed. Investigate false-negative cases.
- [RECOMMENDATION] Accuracy may not represent real-world performance because the benchmark dataset is imbalanced.